

Bayesian Data Analysis, Part 2

Clay Ford

Spring 2024

Agenda

- ▶ Review some basic concepts from Part 1
- ▶ Implement more complex Bayesian analyses
- ▶ model visualization

Review: Bayesian statistics

- ▶ The objective of a Bayesian analysis is a *posterior distribution* of our parameters of interest
- ▶ To obtain a posterior distribution, we must start with a *prior distribution* and a *likelihood*
- ▶ Using *Markov chain Monte Carlo (MCMC)* we sample from the prior and likelihood to estimate the posterior distributions
- ▶ The `rstanarm` package allows us to perform common Bayesian analyses without having to learn Stan


Review: Example

Examine two brands of rechargeable batteries. How long do they run (in hours) before exhausted?

- ▶ Take 25 samples of each brand and calculate mean time
- ▶ Calculate difference between means: 0.86 hours
- ▶ Appears one brand is superior

Review: a prior distribution

Perhaps we're not sure which brand is better. A possible prior distribution might look like this, a $N(0, 3)$ distribution.



Review: Implementation

```
library(rstanarm)
bmod <- stan_glm(y ~ grp, data = bat,
                 family = gaussian,
                 prior = normal(0,3))
```


Bayesian uncertainty interval estimated from the posterior distribution:

```
round(posterior_interval(bmod, prob = 0.95,
                        pars = "grp"),2)
```

```
##      2.5% 97.5%
## grp 0.39  1.33
```

Review: plotting posterior distribution

```
plot(bmod, plotfun = "dens", pars = "grp")
```



Review: Working with posterior distribution

Let's say we want to estimate the probability that the difference in battery life is:

- ▶ greater than 0
- ▶ greater than 1 hour

`as.data.frame()` creates a data frame of *posterior samples*.

```
post <- as.data.frame(bmod)
mean(post$grp > 0)
```


```
## [1] 0.99975
```

```
mean(post$grp > 1)
```

```
## [1] 0.28475
```


Review: graphical posterior predictive checks

```
pp_check(bmod)
```



Moving on to more complex analyses

The previous example was a simple *model* with one predictor. It was essentially the Bayesian approach to what is traditionally analyzed as a t test.

Today's workshop will cover more complicated analyses including:

- ▶ multiple regression
- ▶ regression with interactions
- ▶ binary logistic regression

Multiple regression

Multiple regression, or linear modeling, is the idea that the variability of some numeric variable of interest can be “explained” by a sum of weighted predictors.

Example: patient satisfaction scores at a hospital. Some are high, some are low. Why? Perhaps it has do with their age, anxiety level, and illness severity.

$$\text{satisfaction} = \beta_0 + \beta_1 \text{age} + \beta_2 \text{anxiety} + \beta_3 \text{illness}$$

Where the betas represent some *weight*. Hence the term, *weighted sum*. β_0 is the intercept.

Multiple regression

$$\text{satisfaction} = \beta_0 + \beta_1 \text{age} + \beta_2 \text{anxiety} + \beta_3 \text{illness}$$

This model says, “if I take age, anxiety level and illness severity, multiply each by some weight, and add them up, I’ll get an expected patient satisfaction score.”

The calculated value will be off by some amount. We assume this amount, usually denoted ϵ , is a random draw from a Normal distribution with mean 0 and some unknown standard deviation, σ . This gives us

$$\text{satisfaction} = \beta_0 + \beta_1 \text{age} + \beta_2 \text{anxiety} + \beta_3 \text{illness} + \epsilon$$

Traditional multiple regression means estimating the betas and σ .

Using lm

The traditional approach in R uses the `lm` function.

```
ps <- read.csv("data/patient_satisfaction.csv")  
m <- lm(ps ~ age + illness + anxiety, data = ps)
```

The betas (coefficients/weights) and σ can be viewed with `coef` and `sigma`

```
coef(m)
```

```
## (Intercept)          age      illness      anxiety  
## 158.4912517   -1.1416118   -0.4420043  -13.4701632
```

```
sigma(m)
```

```
## [1] 10.05798
```

Using glm

We can also use `glm` for multiple regression. Note the `family` argument which allows us to specify the error distribution.

```
m2 <- glm(ps ~ age + illness + anxiety, data = ps,  
          family = gaussian)  
coef(m2)
```

```
## (Intercept)          age      illness      anxiety  
## 158.4912517   -1.1416118   -0.4420043  -13.4701632
```

```
sigma(m2)
```

```
## [1] 10.05798
```

The Bayesian approach

As before, instead of estimating parameters, the Bayesian approach is to **estimate the distributions of parameters**.

We propose a prior distribution for the betas and σ , and update those distributions using a likelihood and the data.

Likelihood

Recall our model:

$$\text{satisfaction} = \beta_0 + \beta_1 \text{age} + \beta_2 \text{anxiety} + \beta_3 \text{illness} + \epsilon$$

where $\epsilon \sim N(0, \sigma)$

This implies

$$\text{satisfaction} \sim N(\beta_0 + \beta_1 \text{age} + \beta_2 \text{anxiety} + \beta_3 \text{illness}, \sigma)$$

This is our *likelihood*: A normal, or Gaussian, distribution with a *conditional* mean that changes based on age, anxiety and illness.

Where traditional statistics maximizes likelihood, Bayesian statistics multiplies the likelihood by the prior to get a posterior distribution.

Using stan_glm

The `stan_glm` function uses the same syntax as `glm` and provides *weakly informative* default prior distributions.¹

```
bm <- stan_glm(ps ~ age + illness + anxiety,  
              data = ps,  
              family = gaussian)
```

Instead of point estimates for the betas and σ , we get posterior distributions.

¹On startup, `rstanarm` states “Default priors may change, so it’s safest to specify priors, even if equivalent to the defaults.”

Using stan_glm with explicit priors

Use the prior arguments to specify priors. Below are the default priors. The normal and exponential functions are from rstanarm.

```
bm <- stan_glm(ps ~ age + illness + anxiety,  
              data = ps,  
              family = gaussian,  
              prior_intercept = normal(mean(ps$ps),10),  
              prior = normal(c(0,0,0),c(2.5,2.5,2.5)),  
              prior_aux = exponential(1))
```

Evaluating and exploring the model

We proceed the same way as before to evaluate and explore the model.

```
plot(bm, plotfun = "dens")  
posterior_interval(bm)  
summary(bm)  
pp_check(bm)
```

What about residuals?

In traditional statistics residuals are staple in assessing model fit and diagnostics.

Residuals = Observed response - Predicted response

A traditional model has one set of residuals. A Bayesian model fit using `rstanarm` defaults will have 4000 sets of residuals.

Fortunately the `pp_check()` function provides options for working with Bayesian residuals.

Let's go to R!

Models with interactions

In our previous model, we assumed the predictor effects were simply additive. For example, it didn't matter how ill you were, the effect of age was always the same.

$$\text{satisfaction} = \beta_0 + \beta_1 \text{age} + \beta_2 \text{anxiety} + \beta_3 \text{illness}$$

But we may have reason to believe the effects of age and illness interact. Perhaps the older you are, the effect of illness on patient satisfaction decreases.

One way to describe this interaction is to add the product of age and illness to the model:

$$\text{satisfaction} = \beta_0 + \beta_1 \text{age} + \beta_2 \text{anxiety} + \beta_3 \text{illness} + \beta_4 \text{age} \times \text{illness}$$

Specifying interactions

Use a colon to specify interactions in the model syntax.

```
bm2 <- stan_glm(ps ~ age + illness + anxiety +  
               age:illness,  
               data = ps,  
               family = gaussian)
```

Or use the asterisk as a shortcut: `age * illness = age + illness + age:illness`

The modeling result

Once again the target of the Bayesian model is the collection of posterior distributions on the model weights, or coefficients.

```
posterior_interval(bm2)
```

##		5%	95%
## (Intercept)	26.20181467	251.08202193	
## age	-3.42330550	2.16140201	
## illness	-2.33270503	2.20467749	
## anxiety	-25.84181972	-1.29440540	
## age:illness	-0.06550184	0.04424613	
## sigma	8.65894466	12.60598498	

The interaction appears to be small and we're uncertain whether it's positive or negative.

The modeling result

The `coef` function returns the medians of the posterior distributions of the coefficients.²

```
as.matrix(coef(bm2))
```

```
##                [,1]
## (Intercept) 138.552924803
## age         -0.661732680
## illness     -0.022662797
## anxiety     -13.660200408
## age:illness  -0.009611631
```

²Using `as.matrix` to force the coefficients into a column so they will fit on the slide.

Visualizing interactions

Even if we had good evidence that the coefficient for an interaction was large and in a certain direction (positive or negative), the coefficient can be hard to interpret.

To aid in interpretation we can use *effect plots* to help us visualize our models.

The basic idea is to generate predictions for various combinations of predictors and plot the result.

Using ggeffects to create effect plots


The `ggeffects` package provides methods for easily creating effect plots for models created with `rstanarm`.

The basic syntax to quickly generate an effect plot for our interaction:

```
library(ggeffects)
ggpredict(bm2, terms = c("age", "illness")) |>
  plot()
```

The following plot shows that the interaction between age and illness is small and virtually non-existent.

An effect plot for the age:illness interaction



Customizing the effect plot

By default `ggpredict` will pick some values for the 2nd term. We can specify values if we like as follows:


```
ggpredict(bm2, terms = c("age", "illness[45,50,55]")) |>  
  plot()
```

If we want illness on the x-axis:

```
ggpredict(bm2, terms = c("illness", "age[30,40,50]")) |>  
  plot()
```

An effect plot for the age:illness interaction

```
ggpredict(bm2, terms = c("illness", "age[30,40,50]")) |>  
plot()
```



Interpreting interactions in effect plots


On the previous slide, the effect of illness on patient satisfaction appeared to be the same regardless of age. The plotted lines were approximately parallel. This indicates a small or negligible interaction.

If the plotted lines had vastly different trajectories such that they crossed or grew further apart, then we would have evidence of an interaction.

Effect plots for main effects

Effect plots are useful for main effects as well (ie, predictors not involved in an interaction)

```
ggpredict(bm2, terms = "anxiety") |> plot()
```




The 95% credibility ribbon is for the *mean* response value.

Effect plots for main effects

Set `ppd = TRUE` to get a prediction interval.

```
ggpredict(bm2, terms = "anxiety", ppd = TRUE) |> plot()
```



The 95% credibility ribbon is for the *predicted* response value. Let's go to R!

Logistic regression

Logistic regression models the probability of a binary outcome. The dependent variable is often of the form 0/1, failure/success or no/yes.

Like multiple regression, we model probability as a weighted sum of predictors.

Unlike multiple regression, there is no σ and the weighted sum of predictors are embedded in the *logistic* function:

$$P(Y = 1) = \frac{1}{1 + \exp(-X\beta)}$$

where $X\beta = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k$

This ensures our model always returns values between 0 and 1.

The logit transformation

We can express the logistic regression model as a simple weighted sum of predictors by using the *logit* transformation:

$$\log \left(\frac{P(Y = 1)}{1 - P(Y = 1)} \right) = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$$

In this transformation, the response and coefficients are on the *log odds* scale.

This is the form logistic regression takes when performed in R (or any other program).

Likelihood

Since we're modeling the probability of an event happening, our response variable has a binomial distribution with $n = 1$:

$$Y \sim B\left(n = 1, p = \frac{1}{1 + \exp(-X\beta)}\right)$$

While traditional statistics maximizes this likelihood, Bayesian statistics multiplies the likelihood by the prior to get a posterior distribution.

Logistic regression example

A clinical trial investigates a new treatment for rheumatoid arthritis. Model probability of seeing improvement in condition based on treatment, sex, and age.

```
head(arthritis)
```

##	Treatment	Sex	Age	Better
## 1	Placebo	Female	54	0
## 2	Treated	Female	69	0
## 3	Treated	Male	27	1
## 4	Treated	Female	62	1
## 5	Placebo	Male	44	0
## 6	Treated	Male	70	1

$$\log \left(\frac{P(\text{Better} = 1)}{1 - P(\text{Better} = 1)} \right) = \beta_0 + \beta_1 \text{Trt} + \beta_2 \text{Sex} + \beta_3 \text{Age}$$

Fitting the model

The traditional method uses `glm`. Notice we set `family = binomial` since our response is binary.

```
glm1 <- glm(Better ~ Treatmnt + Sex + Age,  
            data = arthritis,  
            family = binomial)
```

The `rstanarm` specification is virtually identical, except we use `stan_glm`

```
bglm1 <- stan_glm(Better ~ Treatment + Sex + Age,  
                 data = arthritis,  
                 family = binomial)
```

The default coefficient priors are the same as those used when performing multiple regression. The intercept has location = 0.

Interpreting the coefficients

Recall Bayesian modeling does not return point estimates for the coefficients (the betas) but rather distributions. To get a coefficient value, we typically take the median or the mean of the posterior distribution.

The `coef` function returns the median values.

```
coef(bg1m1)
```

##	(Intercept)	TreatmentTreated	SexMale
##	-3.11556378	1.74090082	-1.45083163

Exponentiating the coefficient value returns an odds ratio.

The odds that `Better = 1` is about 6 times higher for the Treated group: $\exp(1.74) \approx 5.7$

Using effect plots with logistic regression models

An effect plot can help communicate a model in terms of probability.

```
ggpredict(bglm1, terms = "Treatment") |> plot()
```



Let's go to R!

Some journal articles using `rstanarm`

Espe, M. et al. (2016) Yield gap analysis of US rice production systems shows opportunities for improvement. *Field Crops Research*, 196:276-283.

Kubrak, O. et al. (2017) Adaptation to fluctuating environments in a selection experiment with *Drosophila melanogaster*. *Ecology and Evolution*, 7:3796-3807.

Herzog, S. et al. (2017) Sun Protection Factor Communication of Sunscreen Effectiveness. *JAMA Dermatology*, 153(3):348-350.

Kovic, M. and Hänsli, N. (2017) The impact of political cleavages, religiosity, and values on attitudes towards nonprofit organizations. *Social Sciences*, 7(1), 2.

References

McElreath, M. (2016). *Statistical Rethinking*. CRC Press. Boca Raton.

Muth, C., Oravecz, Z., & Gabry, J. (2018). User-friendly Bayesian regression modeling: A tutorial with rstanarm and shinystan. *The Quantitative Methods for Psychology*, 14(2), 99-119.

Nicenboim, B., Schad, D. , and Vasishth, S. (2021) An Introduction to Bayesian Data Analysis for Cognitive Science.

<https://vasishth.github.io/bayescogsci/book/>

rstanarm web site: <http://mc-stan.org/rstanarm/>